

Mingyu Sung

AI Research Engineer — Efficient GenAI Inference (Quantization · Caching · Acceleration)

Nota AI, Platform Team (Quantizer) · South Korea

AI Research Engineer at Nota AI (Platform Team, Quantizer part). Ph.D. in Artificial Intelligence (KNU, 2026). Works on efficient generative-model inference: model quantization, diffusion/LLM cache acceleration, CUDA kernel optimization, split computing, and training-free KV-cache compression. Core contributor to the Nunchaku CUDA acceleration engine (ICLR 2025 Spotlight); 15 merged PRs across Nunchaku, ComfyUI-nunchaku, ByteDance DreamO, and Hugging Face Transformers.

EXPERIENCE

AI Research Engineer — Platform Team, Quantizer Part

2026.02 – Present

Nota AI

- LLM/VLM quantization platform: GPTQ, AWQ, SmoothQuant, QuaRot/SpinQuant pipelines; NVFP4 and GGUF export paths; MoE all-expert calibration; KV-cache quantization (verifiable via merged internal PR record).
- Model enablement & evaluation: EXAONE 4.0 / MoE, Qwen3 / 3.5(-VL), Phi-3 (LongRoPE), gpt-oss, embedding/reranker and video-diffusion (Wan, Cosmos) evaluation harnesses.
- K-EXAONE 236B (MoE) optimization for FuriosaAI's data-center NPU, with LG AI Research — ~71% model-size reduction at ~99.2% accuracy retention (GPQA 79.80, IFBench 68.98, AIME25 88.57); announced 2026-06-30.

EDUCATION

Ph.D., Artificial Intelligence

2021.08 – 2026.02

Kyungpook National University (KNU) · Thesis: Fast and Trustworthy Super-Resolution with VLM-Guided Diffusion Model

M.S., Computer Science

2019.09 – 2021.08

Kyungpook National University (KNU) · Thesis: Probabilistic Classification Method of Spiking Neural Network Based on Multi-Labeling of Neurons

B.S., Computer Science

2013.03 – 2019.08

Kyungpook National University (KNU)

RESEARCH

StreamPRS

2026

Single-context-prefill approximation of KVzip's reconstruction-based KV-cache eviction; contributions framed as Factorization / Method / System. (Scope-limited per claims ledger — do not overstate.)

OPEN SOURCE

Nunchaku — Core Contributor

2025 – 2026-02

nunchux-ai/nunchaku · 9 merged PRs · CUDA acceleration engine for 4-bit neural networks (ICLR 2025 Spotlight, 3.9k+ stars)

- Cache optimizations — V2 FBCaching (#621), double FB cache, TeaCache batch processing (#601) — for 2–5x additional speedup
- IP-Adapter (XLabs flux-ip-adapter-v2) support (#418)
- Fixes: LoRA key mismatch (#557), ControlNet (#360, #452), Sana (#380), offload segfault (#440)
- Forward-pass tensor-op cleanup (#491)

ComfyUI-nunchaku — Contributor

2025

nunchux-ai/ComfyUI-nunchaku · 2 merged PRs · ComfyUI integration for the Nunchaku engine

- IP-Adapter support for ComfyUI (#305)
- Cache-mechanism file separation (#474)

DreamO — Contributor

2025

bytedance/DreamO · 1 merged PRs · ByteDance's unified image-customization framework

- Dynamic FBCache / DoubleFBCache support for Nunchaku engine integration (#104)

Transformers — Contributor

2026

huggingface/transformers · 3 merged PRs · Hugging Face Transformers library

- Fixed float16 overflow in Gemma4 vision pooler (#46277)
- Fixed Gemma sliding_window being halved on every config save/reload, a silent accuracy regression for EmbeddingGemma (#47940)
- Fixed _init_weights reading a quantized child's weight, which broke loading CLIP-family quantized checkpoints (#47921)

PUBLICATIONS

[J1] H2-Cache: A Novel Hierarchical Dual-Stage Cache for High-Performance Acceleration of Generative Diffusion Models

M. Sung, I.-M. Kim, S. Yun, J.-M. Kang · IEEE Open Journal of the Computer Society · 2025

- [J2] **A Novel VLM-Guided Diffusion Model for Remote Sensing Image Super-Resolution**
M. Sung, M.-G. Gong, S.-J. Ham, I.-M. Kim, S. Yun, J.-M. Kang · IEEE Geoscience and Remote Sensing Letters · 2025
- [J3] **DeCo-MeSC: Deep Compression-Based Memory-Constrained Split Computing Framework for Cooperative Inference of Neural Network**
M. Sung, V. Palakonda, I.-M. Kim, S. Yun, J.-M. Kang · IEEE Transactions on Vehicular Technology · 2025
- [J4] **Generative Diffusion Model-Based Deep Learning Framework for Remaining Useful Life Prediction**
S. Ha, M. Sung, F. Saeed, S. Yun, I.-M. Kim, J.-M. Kang · IEEE Internet of Things Journal · 2025
- [J5] **Entropy-based sampling for efficient training of deep learning on CNC machining dataset**
M. Sung, C. Park, S. Ha, M. Ha, H. Lee, J. Kim, J.-M. Kang · Electronics Letters · 2024
- [J6] **Probabilistic Classification Method of Spiking Neural Network Based on Multi-Labeling of Neurons**
M. Sung, J. Kim, J.-M. Kang · Mathematics · 2023
- [J7] **Training and Inference using Approximate Floating-Point Arithmetic for Energy Efficient Spiking Neural Network Processors**
M. Kwak, J. Lee, H. Seo, M. Sung, Y. Kim · ICEIC · 2021
- [J8] **Training spiking neural networks with an adaptive leaky integrate-and-fire neuron**
M. Sung, Y. Kim · IEEE ICCE-Asia · 2020

PREPRINTS

- [P1] **GLYPH-SR: Can We Achieve Both High-Quality Image Super-Resolution and High-Fidelity Text Recovery via VLM-guided Latent Diffusion Model?**
arXiv:2510.26339 · 2025
- [P2] **Memory- and Latency-Constrained Inference of Large Language Models via Adaptive Split Computing**
arXiv:2511.04002 · 2025
- [P3] **No Pose Estimation? No Problem: Pose-Agnostic and Instance-Aware Test-Time Adaptation for Monocular Depth Estimation**
arXiv:2511.05055 · 2025
- [P4] **Why Should the Server Do It All?: A Scalable, Versatile, and Model-Agnostic Framework for Server-Light DNN Inference over Massively Distributed Clients via Training-Free Intermediate Feature Compression**
arXiv:2511.11608 · 2025
- [P5] **Range Asymmetric Numeral Systems-Based Lightweight Intermediate Feature Compression for Split Computing of Deep Neural Networks**
arXiv:2511.11664 · 2025

SELECTED PROJECTS

K-EXAONE 236B MoE Optimization (Nota AI × LG AI Research × FuriosaAI) [↗](#) Quantization Engineer · 2026
Optimized the 236B-parameter MoE model for FuriosaAI's data-center NPU via targeted precision analysis of degradation-prone sections — ~71% size reduction with ~99.2% accuracy retention (GPQA 79.80 / IFBench 68.98 / AIME25 88.57).

LLM Inference Optimization (Samsung Challenge) Lead Developer · 2024
Optimized pipeline for Phi-3 using Torch-TensorRT conversion and memory-aware dynamic batching to accelerate LLM inference.

Advanced CNC Machine Tool Diagnosis & Prognosis (KERI) Lead Developer · 2023
Deep learning models for tool wear classification, anomaly detection, and RUL prognosis.

AI-based Intelligent CCTV (ABB Project) AI Optimization Engineer · 2023
Optimized an LSTR behavior-detection model for edge devices via Knowledge Distillation (TransKD) and Deep Compression — 25x model compression.

Global Basic Research Laboratory (NRF Project) Participating Researcher · 2022
Deep learning-based channel estimation (CNN denoising, self-supervised learning) for IRS-aided communication systems.

ML-based Sensor Data Analysis (JS System Project) Lead Developer · 2022
ML-based sensor feature-importance extraction and model performance optimization.

AWARDS

Best Paper Award · KNU-EE Research Congress 2025
A Novel VLM-Guided Diffusion Model for Remote Sensing Image Super-Resolution

SKILLS

Research	model quantization, KV-cache compression, diffusion/LLM inference acceleration, split computing, super-resolution, long-context evaluation
Programming	Python, C++, CUDA C, MATLAB
Frameworks & Tools	PyTorch, HuggingFace Transformers, TensorRT, Triton, flash-attn, Docker, Git, LaTeX
Methodology	preregistration, claims ledger, freeze discipline
Languages	Korean (native) · English (business fluent)